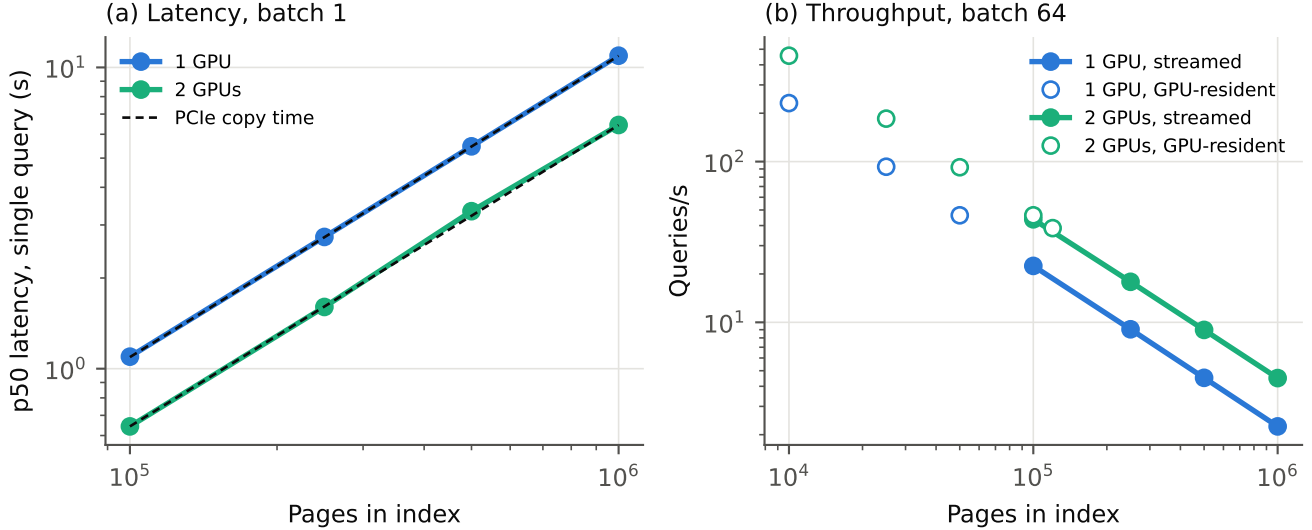
(a) Latency, batch 1
p50 latency, single query (s)
10^1
10^0
1 GPU
2 GPUs
PCIe copy time
10^5
10^6
Pages in index
(b) Throughput, batch 64
Queries/s
10^2
10^1
1 GPU, streamed
1 GPU, GPU-resident
2 GPUs, streamed
2 GPUs, GPU-resident
10^4
10^5
10^6
Pages in index