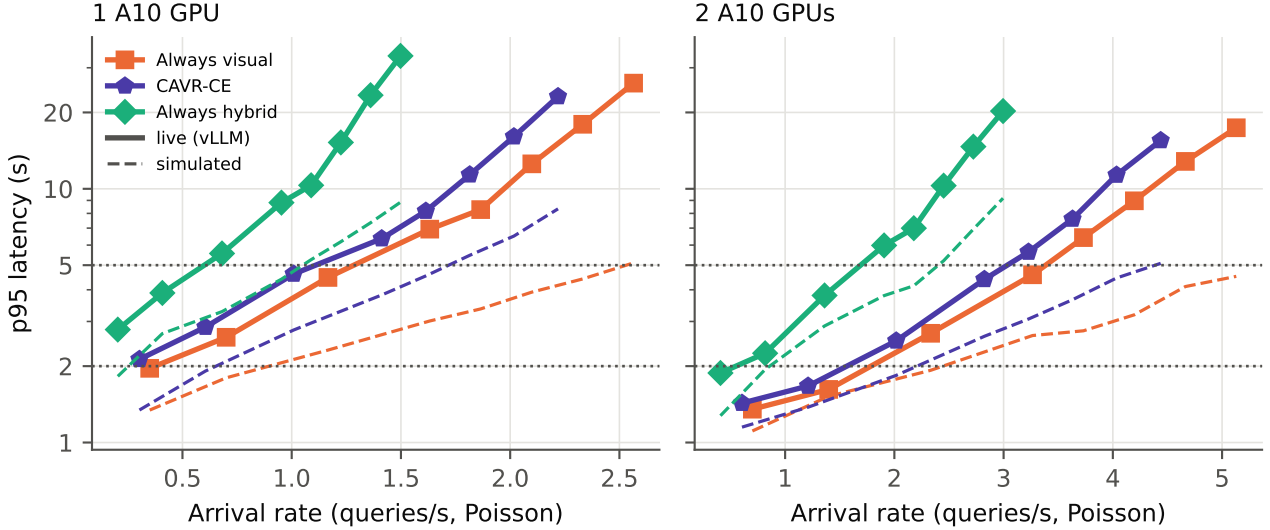
1 A10 GPU
2 A10 GPUs
Always visual
CAVR-CE
Always hybrid
live (vLLM)
simulated
p95 latency (s)
20
10
5
2
1
0.5
1.0
1.5
2.0
2.5
1
2
3
4
5
Arrival rate (queries/s, Poisson)
Arrival rate (queries/s, Poisson)