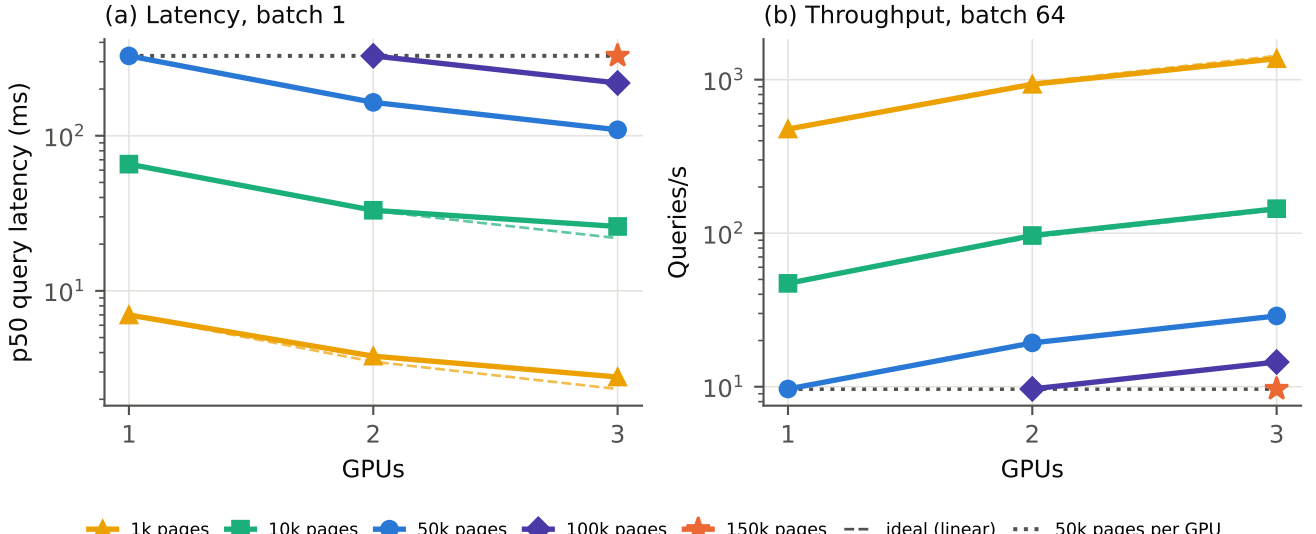
(a) Latency, batch 1
p50 query latency (ms)
10^2
10^1
1
2
3
GPUs
(b) Throughput, batch 64
Queries/s
10^3
10^2
10^1
1
2
3
GPUs
1k pages
10k pages
50k pages
100k pages
150k pages
ideal (linear)
50k pages per GPU