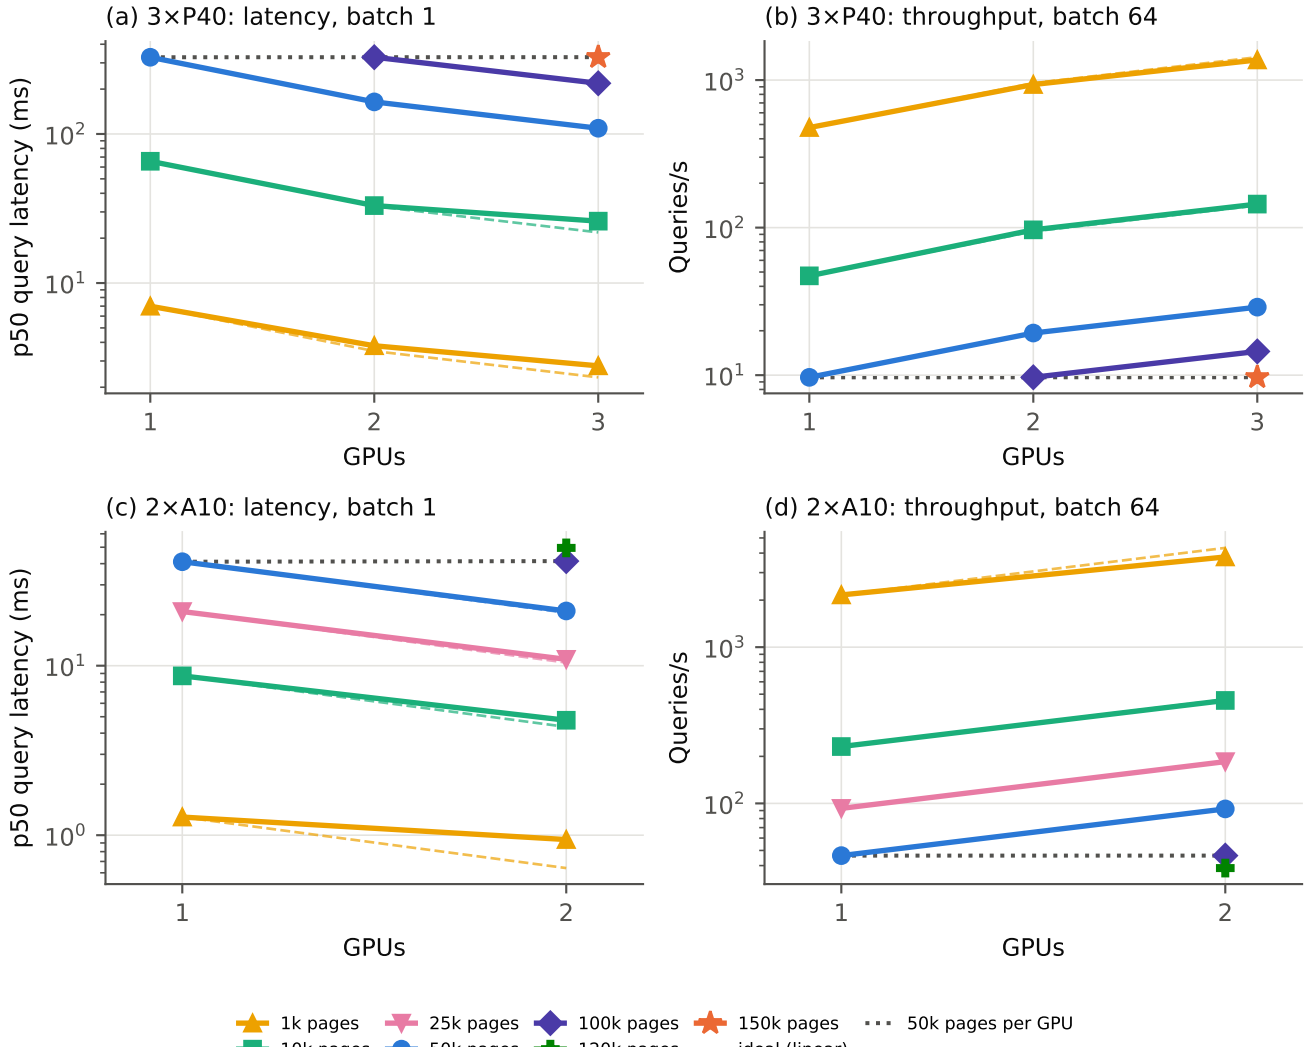
(a) 3×P40: latency, batch 1
p50 query latency (ms)
10^2
10^1
1
2
3
GPUs
(b) 3×P40: throughput, batch 64
Queries/s
10^3
10^2
10^1
1
2
3
GPUs
(c) 2×A10: latency, batch 1
p50 query latency (ms)
10^1
10^0
1
2
GPUs
(d) 2×A10: throughput, batch 64
Queries/s
10^3
10^2
1
2
GPUs
1k pages
10k pages
25k pages
50k pages
100k pages
120k pages
150k pages
50k pages per GPU
ideal (linear)